

# ACO Laboratory Assignment

## Lab 5: Geometric (data) decomposition using implicit tasks: heat diffusion equation

---

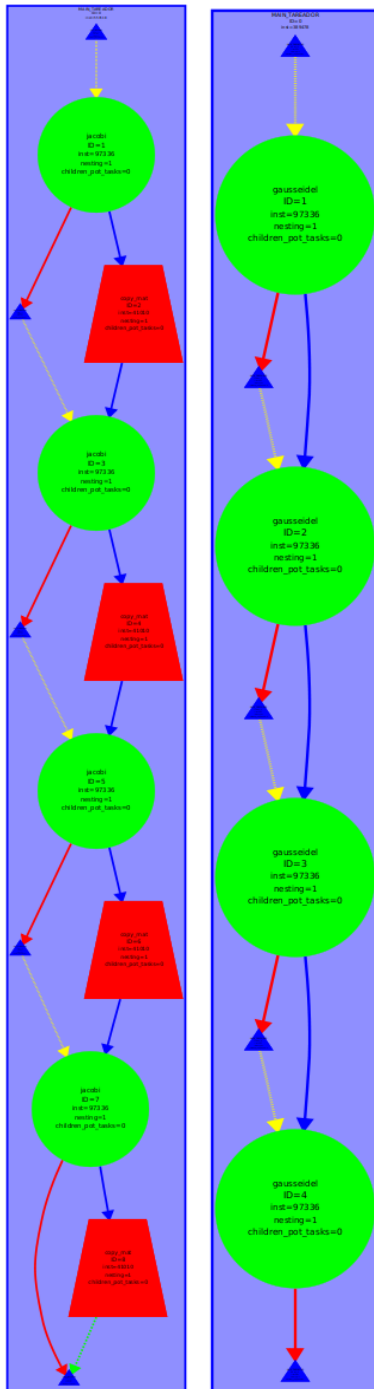
Oriol Amate Sabat i Adrian Garcia Campillo

# Index

<b>Index.....</b>	<b>2</b>
<b>2.Sequential heat diffusion program and analysis with Tareador.....</b>	<b>3</b>
Podem paralelitzar copy_mat y solve.....	3
<b>3.Parallelisation of the heat equation solvers.....</b>	<b>10</b>
3.1 Jacobi solver.....	10
3.1.2 Overall Analysis.....	11
3.1.3 Detailed Analysis.....	12
3.1.4 Optimization.....	13
Overall Analysis of the Optimized Code.....	13
Detailed Analysis of the Optimized Code.....	14
3.2 Gauss–Seidel solver.....	15
3.2.1 First Implementation.....	15
3.2.2 Overall Analysis.....	15
3.2.4 Finding the appropriate value for the number of blocks.....	15

## 2. Sequential heat diffusion program and analysis with Tareador

Include the task dependency graph shown by Tareador. Is there any parallelism that can be exploited at this granularity level?



Podem parallelitzar `copy_mat` y `solve`.

Include the excerpt of the code that you have modified in order to specify one task per block.


```

12 // Function to copy one matrix into another
13 void copy_mat (double *u, double *v, unsigned sizex, unsigned sizey) {
14
15     int nblocksx=4;
16     int nblocksy=4;
17
18     for (int blockx=0; blockx<nblocksx; ++blockx) {
19         int i_start = lowerb(blockx, nblocksx, sizex);
20         int i_end = upperb(blockx, nblocksx, sizex);
21         for (int blocky=0; blocky<nblocksy; ++blocky) {
22             int j_start = lowerb(blocky, nblocksy, sizey);
23             int j_end = upperb(blocky, nblocksy, sizey);
24             tareador_start_task("jacobi");
25             for (int i=max(1, i_start); i<=min(sizex-2, i_end); i++)
26                 for (int j=max(1, j_start); j<=min(sizey-2, j_end); j++)
27                     v[i*sizey+j] = u[i*sizey+j];
28             tareador_end_task("jacobi");
29         }
30     }
31 }
32
33 // 2D-blocked solver: one iteration step
34 double solve (double *u, double *unew, unsigned sizex, unsigned sizey) {
35     double tmp, diff, sum=0.0;
36
37     int nblocksx=4;
38     int nblocksy=4;
39
40     //tareador_disable_object(&sum);
41     for (int blockx=0; blockx<nblocksx; ++blockx) {
42         int i_start = lowerb(blockx, nblocksx, sizex);
43         int i_end = upperb(blockx, nblocksx, sizex);
44         for (int blocky=0; blocky<nblocksy; ++blocky) {
45             int j_start = lowerb(blocky, nblocksy, sizey);
46             int j_end = upperb(blocky, nblocksy, sizey);
47             tareador_start_task("block");
48             for (int i=max(1, i_start); i<=min(sizex-2, i_end); i++) {
49                 for (int j=max(1, j_start); j<=min(sizey-2, j_end); j++) {
50                     tmp = 0.25 * ( u[ i*sizey      + (j-1) ] + // left
51                                 u[ i*sizey      + (j+1) ] + // right
52                                 u[ (i-1)*sizey + j      ] + // top
53                                 u[ (i+1)*sizey + j      ] ); // bottom
54                     diff = tmp - u[i*sizey+ j];
55                     sum += diff * diff;
56                     unew[i*sizey+j] = tmp;
57                 }
58             }
59             tareador_end_task("block");
60         }
61     }
62     //tareador_enable_object(&sum);
63

```

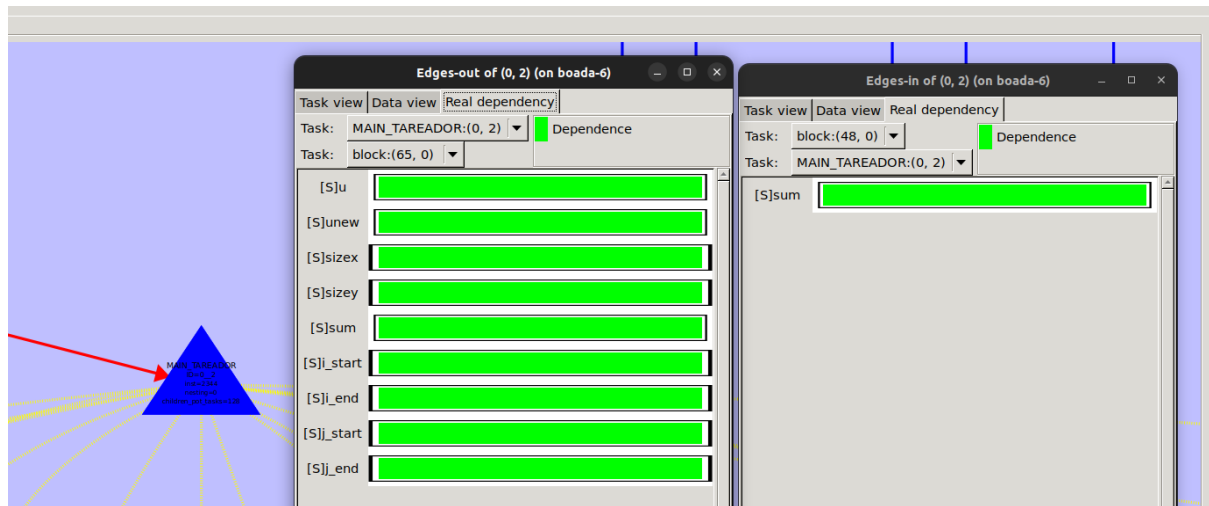
Jacobi

Gauss

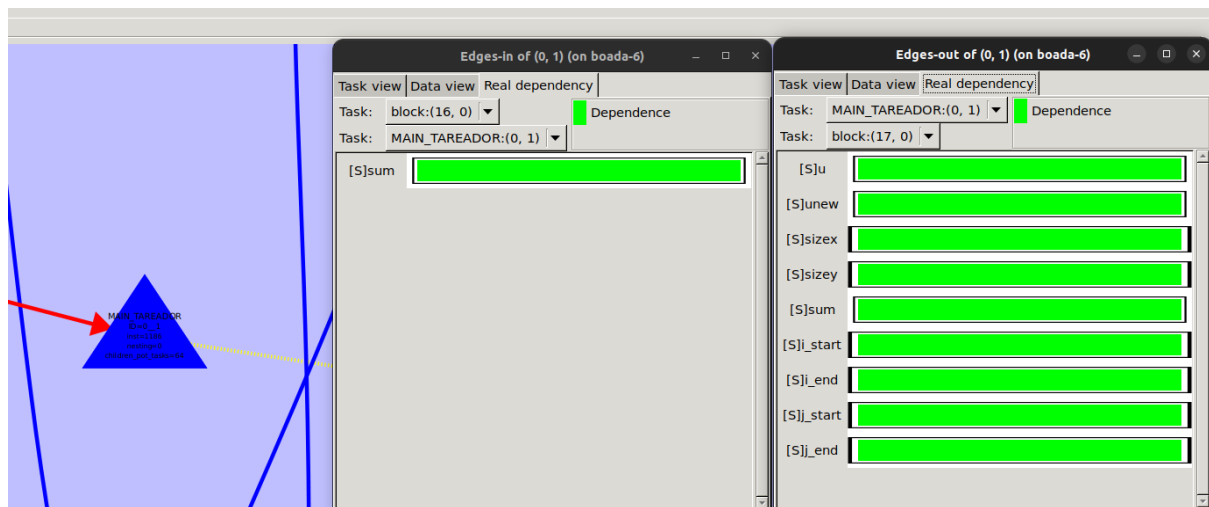


Which variable is causing the serialisation of all the tasks? Use the Dataview option in Tareador to identify it.

Jacobi



Gauss




In order to emulate the effect of protecting the dependences caused by this variable, you can use the tareador disable object and tareador enable object calls, already introduced in the code as comments. With these calls you are telling to Tareador to filter the dependences caused by the variable indicated as object. Uncomment them, recompile and execute.

```


12 // Function to copy one matrix into another
13 void copy_mat (double *u, double *v, unsigned sizex, unsigned sizey) {
14
15     int nblocks_i=4;
16     int nblocks_j=4;
17
18     for (int block_i=0; block_i<nblocks_i; ++block_i) {
19         int i_start = lowerb(block_i, nblocks_i, sizex);
20         int i_end = upperb(block_i, nblocks_i, sizex);
21         for (int block_j=0; block_j<nblocks_j; ++block_j) {
22             int j_start = lowerb(block_j, nblocks_j, sizey);
23             int j_end = upperb(block_j, nblocks_j, sizey);
24             tareador_start_task("jacobi");
25             for (int i=max(1, i_start); i<=min(sizex-2, i_end); i++)
26                 for (int j=max(1, j_start); j<=min(sizey-2, j_end); j++)
27                     v[i*sizey+j] = u[i*sizey+j];
28             tareador_end_task("jacobi");
29         }
30     }
31 }
32
33 // 2D-blocked solver: one iteration step
34 double solve (double *u, double *unew, unsigned sizex, unsigned sizey) {
35     double tmp, diff, sum=0.0;
36
37     int nblocks_i=4;
38     int nblocks_j=4;
39
40     tareador_disable_object(&sum);
41     for (int block_i=0; block_i<nblocks_i; ++block_i) {
42         int i_start = lowerb(block_i, nblocks_i, sizex);
43         int i_end = upperb(block_i, nblocks_i, sizex);
44         for (int block_j=0; block_j<nblocks_j; ++block_j) {
45             int j_start = lowerb(block_j, nblocks_j, sizey);
46             int j_end = upperb(block_j, nblocks_j, sizey);
47             tareador_start_task("block");
48             for (int i=max(1, i_start); i<=min(sizex-2, i_end); i++) {
49                 for (int j=max(1, j_start); j<=min(sizey-2, j_end); j++) {
50                     tmp = 0.25 * ( u[ i*sizey      + (j-1) ] + // left
51                                 u[ i*sizey      + (j+1) ] + // right
52                                 u[ (i-1)*sizey + j      ] + // top
53                                 u[ (i+1)*sizey + j      ] ); // bottom
54                     diff = tmp - u[i*sizey+ j];
55                     sum += diff * diff;
56                     unew[i*sizey+j] = tmp;
57                 }
58             }
59             tareador_end_task("block");
60         }
61     }
62     tareador_enable_object(&sum);
63
64     return sum;
65 }

```

JACOBI



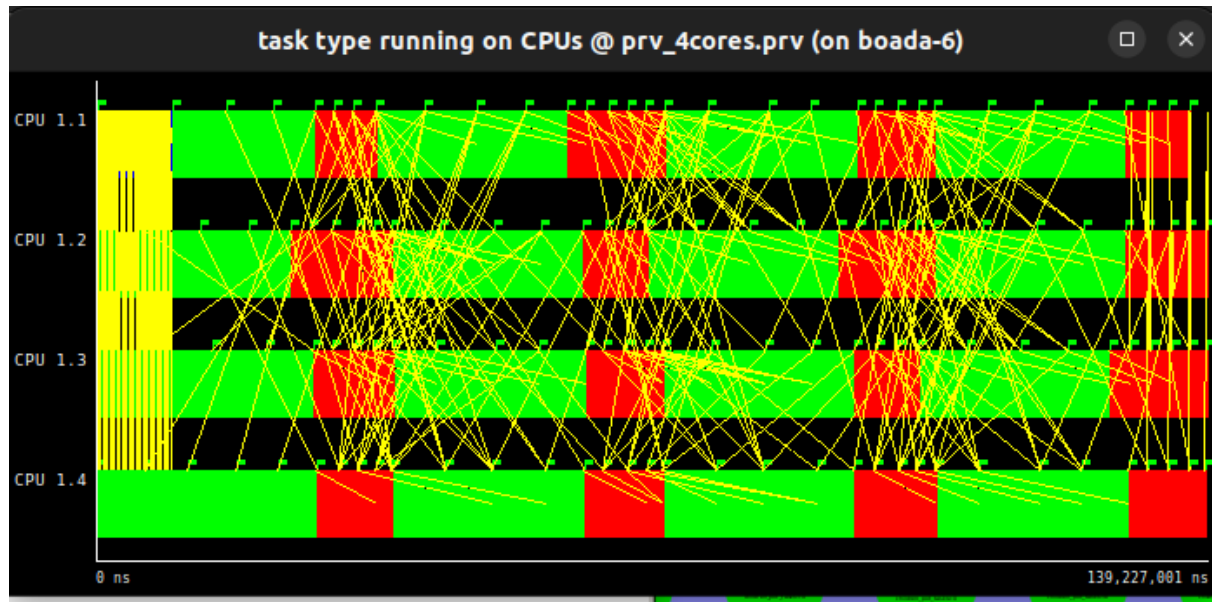
GAUSS






Simulate the execution of both solvers when using 4 threads. Is there any other part of the code that can also be parallelised (take into account this question for the parallel version!)? If so, modify again the instrumentation to parallelise it.

JACOBI



GAUSS



## 3.Parallelisation of the heat equation solvers

### 3.1 Jacobi solver

```
1
2 // 2D-blocked solver: one iteration step
3 double solve (double *u, double *unew, unsigned sizex, unsigned sizey) {
4     double tmp, diff, sum=0.0;
5
6     if (u==unew) return solve_gauss(u,sizex,sizey);
7
8     int nblocksi=omp_get_max_threads();
9     int nblocksj=1;
10
11     #pragma omp parallel private (diff,tmp) reduction(+:sum)
12     // complete data sharing constructs here
13     {
14         int blocki = omp_get_thread_num();
15         int i_start = lowerb(blocki, nblocksi, sizex);
16         int i_end = upperb(blocki, nblocksi, sizex);
17         for (int blockj=0; blockj<nblocksj; ++blockj) {
18             int j_start = lowerb(blockj, nblocksj, sizey);
19             int j_end = upperb(blockj, nblocksj, sizey);
20             for (int i=max(1, i_start); i<=min(sizex-2, i_end); i++) {
21                 for (int j=max(1, j_start); j<=min(sizey-2, j_end); j++) {
22                     tmp = 0.25 * ( u[ i*sizey + (j-1) ] + // left
23                                 u[ i*sizey + (j+1) ] + // right
24                                 u[ (i-1)*sizey + j ] + // top
25                                 u[ (i+1)*sizey + j ] ); // bottom
26                     diff = tmp - u[i*sizey+ j];
27                     sum += diff * diff;
28                     unew[i*sizey+j] = tmp;
29                 }
30             }
31         }
32     }
33
34     return sum;
35 }
```

```

paco1106@boada-6:~/lab5$ cat heat-omp-jacobi-1-boada-11.txt
Iterations      : 25000
Resolution      : 254
Residual        : 0.000050
Solver          : 0 (Jacobi)
Num. Heat sources : 2
  1: (0.00, 0.00) 1.00 2.50
  2: (0.50, 1.00) 1.00 2.50
Time: 2.662
Flops and Flops per second: (11.182 GFlop => 4200.23 MFlop/s)
Convergence to residual=0.000050: 15756 iterations
paco1106@boada-6:~/lab5$ sbatch submit-omp.sh heat-omp 0 8
Submitted batch job 160502
paco1106@boada-6:~/lab5$ ls
fake.h          heat-seq        run-tareador.sh  submit-strong-extrae.sh
heat-gauss.ppm  heat-seq.c      solver-omp.c      submit-strong-omp.sh
heat.h          heat-tareador   solver-seq.c      submit-userparam-omp.sh
heat-jacobi.ppm heat-tareador.c solver-tareador.c  table-generation.sh
heat-omp        logs           submit-omp.sh     tareador_llvm.log
heat-omp.c      Makefile       submit-omp.sh.e160499 test.dat
heat-omp-jacobi-1-boada-11.txt misc.c         submit-omp.sh.e160502 userparam-omp.jgr
heat-omp-jacobi-8-boada-12.txt modelfactor-tables.tex submit-omp.sh.o160499 verbose.log
heat.ppm        run-lite-tareador.sh submit-omp.sh.o160502
paco1106@boada-6:~/lab5$ cat heat-omp-jacobi-8-boada-12.txt
Iterations      : 25000
Resolution      : 254
Residual        : 0.000050
Solver          : 0 (Jacobi)
Num. Heat sources : 2
  1: (0.00, 0.00) 1.00 2.50
  2: (0.50, 1.00) 1.00 2.50
Time: 4.074
Flops and Flops per second: (11.182 GFlop => 2744.45 MFlop/s)
Convergence to residual=0.000050: 15756 iterations
paco1106@boada-6:~/lab5$

```

### 3.1.2 Overall Analysis

Include the modelfactor tables. Is the scalability that is obtained with this initial parallelisation appropriate? Which is the metric reported by mflite.py that you should address first?

Veiem que encara es pot paralelitzar si toquem copy\_mat

Overview of whole program execution metrics				
Number of processors	1	4	8	16
Elapsed time (sec)	3.09	2.16	2.08	2.13
Speedup	1.00	1.43	1.48	1.45
Efficiency	1.00	0.36	0.19	0.09

Table 1: Analysis done on Sun Dec 24 01:51:07 PM CET 2023, paco1106

Overview of the Efficiency metrics in parallel fraction, $\phi=67.34\%$				
Number of processors	1	4	8	16
Global efficiency	99.75%	79.15%	67.11%	37.86%
Parallelization strategy efficiency	99.75%	98.99%	98.07%	97.55%
Load balancing	100.00%	99.93%	99.90%	99.89%
In execution efficiency	99.75%	99.06%	98.16%	97.65%
Scalability for computation tasks	100.00%	79.95%	68.43%	38.81%
IPC scalability	100.00%	84.31%	78.10%	51.71%
Instruction scalability	100.00%	96.09%	94.81%	84.30%
Frequency scalability	100.00%	98.69%	92.41%	89.03%

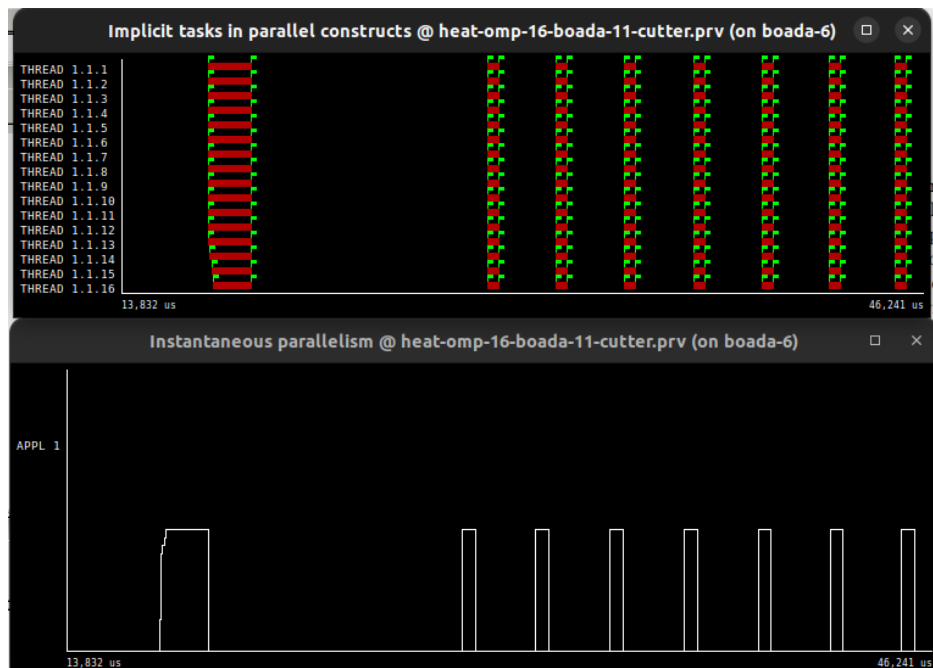
Table 2: Analysis done on Sun Dec 24 01:51:07 PM CET 2023, paco1106

Statistics about explicit tasks in parallel fraction				
Number of processors	1	4	8	16
Number of implicit tasks per thread (average us)	1000.0	1000.0	1000.0	1000.0
Useful duration for implicit tasks (average us)	2076.36	649.23	379.27	334.35
Load balancing for implicit tasks	1.0	1.0	1.0	1.0
Time in synchronization implicit tasks (average us)	0	0	0	0
Time in fork/join implicit tasks (average us)	5.2	0	0	0

Table 3: Analysis done on Sun Dec 24 01:51:07 PM CET 2023, paco1106

### 3.1.3 Detailed Analysis

Include some captures of the window timelines to show the problem. What is the region of the code that is provoking the low value for the parallel fraction in your parallelisation? Hint: Remember the Tareador analysis you did.



Falta paralelitzar copy\_mat.

### 3.1.4 Optimization

Include an excerpt of the code to show the OpenMP annotations you have added to the code.

Parallelized copy\_mat

```
// Function to copy one matrix into another
void copy_mat (double *u, double *v, unsigned sizex, unsigned sizey) {
    int nbblocksi=1;
    int nbblocksj=1;
    #pragma omp parallel
    {
        for (int blocki=0; blocki<nbblocksi; ++blocki) {
            int i_start = lowerb(blocki, nbblocksi, sizex);
            int i_end = upperb(blocki, nbblocksi, sizex);
            for (int blockj=0; blockj<nbblocksj; ++blockj) {
                int j_start = lowerb(blockj, nbblocksj, sizey);
                int j_end = upperb(blockj, nbblocksj, sizey);
                for (int i=max(1, i_start); i<=min(sizex-2, i_end); i++)
                    for (int j=max(1, j_start); j<=min(sizey-2, j_end); j++)
                        v[i*sizex+j] = u[i*sizex+j];
            }
        }
    }
}
```

### Overall Analysis of the Optimized Code

Include the model factors tables and the plot of scalability. Is the execution time reduced?. Have you increased the scalability? Is the parallel fraction larger than before? What is the speedup that you have achieved compared to your first implementation for 16 threads?

Overview of whole program execution metrics				
Number of processors	1	4	8	16
Elapsed time (sec)	3.10	4.71	3.49	3.36
Speedup	1.00	0.66	0.89	0.92
Efficiency	1.00	0.66	0.89	0.92

Table 1: Analysis done on Sun Dec 24 01:41:42 PM CET 2023, paco1106

Overview of the Efficiency metrics in parallel fraction, $\phi=98.89\%$				
Number of processors	1	4	8	16
Global efficiency	99.72%	65.38%	88.50%	92.02%
Parallelization strategy efficiency	99.72%	76.86%	79.93%	82.78%
Load balancing	100.00%	88.02%	83.19%	85.68%
In execution efficiency	99.72%	87.33%	96.07%	96.62%
Scalability for computation tasks	100.00%	85.06%	110.72%	111.16%
IPC scalability	100.00%	31.56%	31.19%	25.26%
Instruction scalability	100.00%	271.99%	383.32%	486.50%
Frequency scalability	100.00%	99.09%	92.62%	90.44%

Table 2: Analysis done on Sun Dec 24 01:41:42 PM CET 2023, paco1106

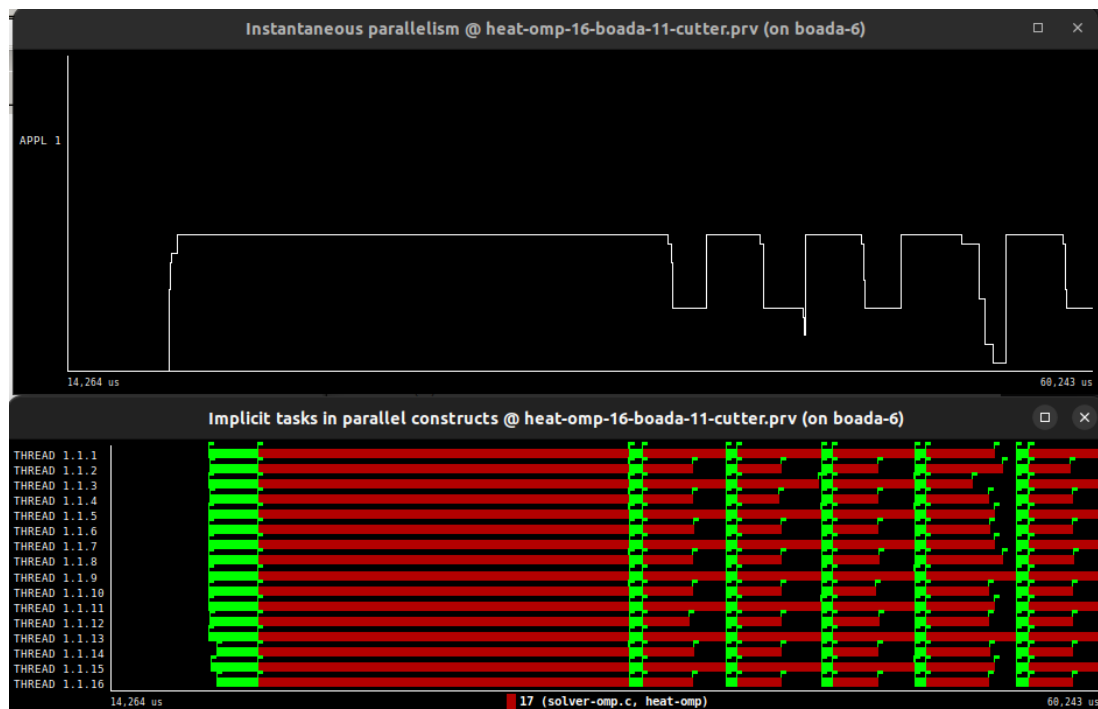
Statistics about explicit tasks in parallel fraction				
Number of processors	1	4	8	16
Number of implicit tasks per thread (average us)	2000.0	2000.0	2000.0	2000.0
Useful duration for implicit tasks (average us)	1529.76	1798.52	1381.61	1376.16
Load balancing for implicit tasks	1.0	0.88	0.83	0.86
Time in synchronization implicit tasks (average us)	0	0	0	0
Time in fork/join implicit tasks (average us)	4.27	0	0	0

Table 3: Analysis done on Sun Dec 24 01:41:42 PM CET 2023, paco1106

Observer que el temps d'execució és una mica mayor, però millora considerablement l'eficiencia.

## Detailed Analysis of the Optimized Code

Include some captures of the window timelines.



## **3.2 Gauss–Seidel solver**

### **3.2.1 First Implementation**

Include an excerpt of the code to show the modifications done.

### **3.2.2 Overall Analysis**

Include the plot obtained with `submit-strong-omp.sh`. Do you observe a linear speedup? Reason your answer. Note: Remember the analysis done with `tareador` with a simulation with 4 threads.

### **3.2.3 Detailed Analysis and Optimization**

Include an excerpt of the code to show the modifications done.

### **3.2.4 Finding the appropriate value for the number of blocks**

Include the plots obtained with `submit-userparam-omp.sh` and `submit-strong-omp.sh` using the best `userparam` and 16 threads. Reason why changing the number of blocks in the `j` dimension changes the ratio between computation and synchronisation. Compare the strong scalability for both cases `nblocksj=20` and `nblocksj=best userparam*nblocksi` (Finding the appropriate value for the number of blocks).